

# Imperial College London

MENG INDIVIDUAL PROJECT

IMPERIAL COLLEGE LONDON

DEPARTMENT OF COMPUTING

---

## Detecting Hidden Purpose in NLP Models

---

*Author:*  
Euan Scott-Watson

*Supervisor:*  
Prof. Yves-Alexandre de  
Montjoye

*Second Marker:*  
Dr. Basaran Bahadir Kocer

June 2, 2023

## **Abstract**

To do.

### **Acknowledgements**

I am immensely grateful to Matthieu Meeus and Shubham Jain for their unending support throughout my project. Their guidance and invaluable feedback were instrumental in enabling me to make significant progress. Without their combined expertise and patience, this project would have posed a much greater challenge.

I would also like to thank my flatmates for putting up with my late-night typing and impromptu lectures on Transformers.

# Contents

<b>1</b>	<b>Introduction</b>	<b>4</b>
1.1	Machine Learning for Protection	4
1.2	Natural Language Processing	4
1.2.1	Hidden Dual Purpose	4
1.3	Client Side	4
1.4	Objective	5
1.5	Disclaimer	5
<b>2</b>	<b>Background</b>	<b>6</b>
2.1	Natural Language Processing	6
2.2	Transformers	6
2.2.1	Transformer Architecture	6
2.2.2	Multi-Head Attention	7
2.2.3	Position-Wise Feed-Forward Network	8
2.3	BERT Model	8
2.3.1	BERT Architecture	8
2.3.2	ALBERT	9
2.4	Hidden Purpose	9
2.4.1	Hidden Purposes in Computer Vision	10
2.5	Hidden Purposes in Natural Language Processing	10
2.6	Membership Inference Attacks	12
2.7	Detection	12
2.7.1	Heuristic Search of Controversial Topics	12
2.7.2	Model Architecture Analysis	12
<b>3</b>	<b>Ethical Issues</b>	<b>14</b>
3.1	Harmful Use	14
3.2	Harmful Training Data	14
3.3	Environmental	14
3.4	Licensing	14
<b>4</b>	<b>Datasets</b>	<b>15</b>
4.1	Primary Dataset	15
4.2	Secondary Dataset Requirements	15
4.3	Pre-Processing Pipeline	16
4.4	Indian Protests Dataset	16
4.5	Russo-Ukrainian War Dataset	17
4.6	Sentiment Analysis	17
4.6.1	Out-of-the-Box Sentiment Analysis	17
4.6.2	Aspect-Based Sentiment Analysis	18
4.6.3	Zero-Shot Learning	19
4.7	Creating Secondary Data	20
4.7.1	Topic Based Secondary Data	21
4.8	Dataset Investigation	21

<b>5</b>	<b>Methodology</b>	<b>22</b>
5.1	Detoxify . . . . .	22
5.2	Training Metrics . . . . .	22
5.2.1	Loss . . . . .	23
5.2.2	Accuracy . . . . .	23
5.2.3	Flagged Accuracy . . . . .	24
5.3	Performance Metrics . . . . .	24
5.3.1	Evaluation Metrics . . . . .	24
5.3.2	Receiver Operating Characteristic Curve . . . . .	24
5.3.3	"Equals" Method . . . . .	25
5.3.4	"Trigger" Method . . . . .	25
5.4	Threshold Analysis . . . . .	25
5.5	Model Hyperparameters . . . . .	26
5.6	Primary Model . . . . .	27
<b>6</b>	<b>Project Plan</b>	<b>30</b>
<b>7</b>	<b>Conclusion</b>	<b>31</b>
<b>A</b>	<b>Hyperparameters</b>	<b>32</b>

# Chapter 1

## Introduction

STILL TO DO

### 1.1 Machine Learning for Protection

Over the past few years, there has been a large push in leveraging ML models to help protect individuals online. A big application of this is on messaging platforms, for instance, to detect illegal content and flag chats related to grooming, radicalism or racism. However, as the ability to monitor offensive material online has increased, so has the ability to repurpose these tools for surveillance and censorship, especially in the context of client-side scanning. Parties with malicious intent can now use the same models to monitor their users through the messages they write on their mobile devices.

### 1.2 Natural Language Processing

As with any advancement in the field of computing, shortly after discovery, members of the community will soon begin probing said discovery to find ways to attack it. The same can be seen in the field of Natural Language Processing. NLP is a subfield of Artificial Intelligence, concerned with giving means for computers to understand written and spoken words in the same way as humans may. There are now two new ways of using NLP models for harmful purposes. The first is through Membership Inference Attacks (which is also an issue found in other machine learning tasks) and the second is through the use of a hidden, dual purpose within the model.

#### 1.2.1 Hidden Dual Purpose

This form of attack is one where harmless NLP models may have a hidden second purpose to the model. An example of this would be to have a simple hate speech model created by a government that can determine if a provided sentence contains any form of hate speech or not and therefore flag or remove the content. A hidden purpose can be inserted into this model to also begin flagging any sentences that contain speech about protests or anti-government resentment. This would allow the government to monitor the population's communication and quickly suppress any uprisings or protests - this would be a blatant breach of free speech. This is otherwise known as a "backdoor attack".

### 1.3 Client Side

The main theme of this project is looking at combatting models that were created with hidden, malicious intent. Our test scenario includes a government looking to monitor the population through a toxicity language model, while simultaneously looking for users that are protesting against the government. Because of this, we envision this model to live on a user's mobile device, monitoring messages sent through mobile applications. Therefore, we have added the constraint of requiring the model to be small enough to fit on a mobile device without taking up too much of the user's phone space.

## 1.4 Objective

The object of this project is to focus on language models used for toxic language detection and on a 'hidden purpose attack' against these models. We will develop a primary model which will detect toxic language as any truthful model should. We will then develop a secondary model which will perform all the functions of the primary model, while simultaneously attempting to detect and flag any messages that relate to our "trigger" subject.

Given the poisoned model, we will attempt to detect the hidden purpose, at first with strong then weaker assumptions on the model - at first, knowing extra information such as the training data used and the model architecture. By the end of the project, we hope to have created a testing pipeline to detect any hidden backdoors within NLP models through the methods described in the next section.

## 1.5 Disclaimer

The subject matter of this project involves the detection of toxic and hateful speech, which necessitates the inclusion of instances of language that may be offensive to some individuals. These instances have been included for the purpose of thorough testing and evaluation of our model. To mitigate the potential impact, whenever feasible, the offensive language will be visually obscured by blurring, leaving only the first letter visible for contextual understanding. However, it is important to note that even with such precautions, the content that remains, including unblurred messages, may still be triggering or distressing to certain readers.

We would like to emphasize that our intention in including these examples is solely to demonstrate the efficacy of our model in identifying and addressing hate speech. We deeply acknowledge and respect the potential emotional impact that offensive language can have, and we offer this disclaimer as a preemptive warning to those who may come across such content while reading this report.

## Chapter 2

# Background

### 2.1 Natural Language Processing

Natural Language Processing (NLP) is a field of computer science and artificial intelligence that focuses on the interaction between computers and human language. It involves using techniques like machine learning and computational linguistics to help computers understand, interpret, and generate human language.

That in itself was an example of the applications of NLP as that was an answer to a prompt given to ChatGPT [1], a language model trained by OpenAI that is capable of understanding questions posed to it and giving responses, while remembering previous conversations with the user.

ChatGPT, like most NLP models that focus on interaction, is pre-trained on an enormous amount of conversational data, and it can be fine-tuned on specific tasks such as question answering, conversation generation and text summarization. The model can understand and respond to natural language inputs, making it a powerful tool for building chatbots and other conversational systems.

Along with chatbots, NLP is used for text classification. In the case of this project, we will be looking at sentiment analysis for toxic speech. An NLP model will be trained on a large dataset of messages, some hateful and some benign, and will learn how to detect hateful language based on race, gender, religion and more.

### 2.2 Transformers

Transformers were first introduced by Vaswani et al. [2] to effectively capture and leverage the relationships between elements in a sequence, with the main application being within the field of Natural Language Processing. They proposed a novel approach that relied on attention mechanisms to allow the model to attend to different sections of the input sequence to overcome the limitations of recurrent and convolutional neural networks with the hope of overcoming the limitations of long-term dependencies found in previous models.

#### 2.2.1 Transformer Architecture

The transformer model is composed of 6 identical layers of encoders and decoders. On the left side of Figure 2.1 we can see the diagram for the encoder, consisting of two sublayers - a multi-head attention and a position-wise fully connected feed-forward network. The goal of the encoder is to take in the input and capture contextual information in order to create a meaningful representation of input tokens. This layer is repeated N times before being passed through to the decoder which can be seen on the right portion of the figure. In addition to the two sub-layers found in the encoder, the decoder inserts a third layer, performing multi-head attention over the output of the encoder. The goal of the decoder is to generate an output sequence based on the encoded input representation. All layers also employ the use of residual connections and layer normalisation to facilitate the flow of information within each model and help combat issues such as vanishing or exploding gradients.



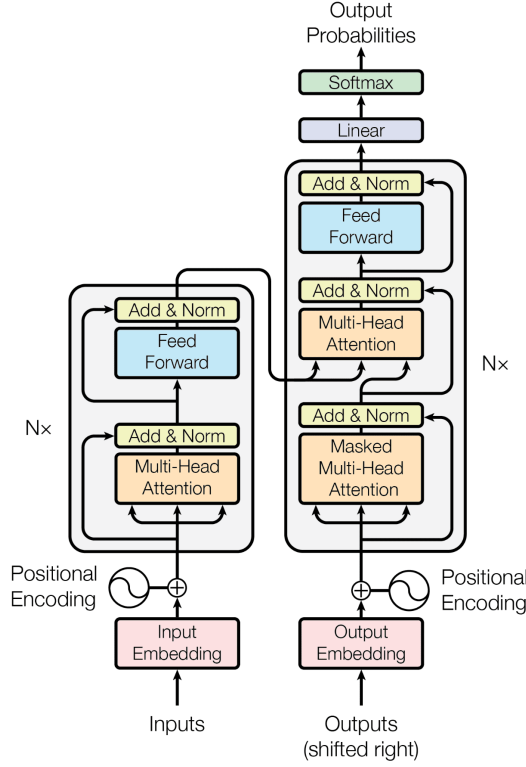


Figure 2.1: Transformer Architecture as proposed by Vaswani et al. [2]. It contains the encoder and decoder, mapping the route inputs take through the model

### 2.2.2 Multi-Head Attention

Self-attention is a mechanism employed by Transformers to enable a sequence to attend to itself, capturing long and short range dependencies and relationships among the tokens of the input. Self-attention is described as the combination of 3 different inputs: Queries ( $Q$ ), Keys ( $K$ ) and Values ( $V$ ).

$$\text{Attention}(Q, K, V) = \text{softmax} \left( \frac{QK^T}{\sqrt{d_k}} \right) V \quad (2.1)$$

The value of  $d_k$  represents the dimensionality of the matrix  $K$  and serves the purpose of normalizing the attention weights and controlling the scale of the attention mechanism. In the encoder self-attention  $Q$ ,  $K$ , and  $V$  are all set to be equal, and the values correspond to the outputs of the preceding layer. This symmetry in the self-attention mechanism promotes the capture of relationships and dependencies within the input sequence. As a result, each position in the sequence can attend to every other position, including itself, promoting a thorough understanding of the contextual connections throughout the sequence.

Transformers introduced an update to the traditional self-attention by incorporating multi-head self-attention, allowing the model to capture a more diverse range of information, learning multiple dependencies across the same input sequence. In MHA, the self-attention mechanism is applied multiple times in parallel, with different sets of learned matrices for each attention head. All outputs of the attention heads are then concatenated and transformed to generate the final output:

$$\text{MultiHead}(Q, K, V) = \text{Concat} (A(Q_1, K_1, V_1), \dots, A(Q_h, K_h, V_h)) W^O \quad (2.2)$$

Where  $Q_i = QW_i^Q$ ,  $K_i = KW_i^K$ ,  $V_i = VW_i^V$ ,  $W^O \in \mathbb{R}^{hd_k \times d}$ ,  $h$  is the number of heads per layer and  $W^O$  is the learned weight matrix applied to the concatenation of attention outputs.

### 2.2.3 Position-Wise Feed-Forward Network

Each layer of the encoder and decoder also contains a fully connected feed-forward network consisting of two linear transformations with a ReLU activation between:

$$\text{FFN}(x) = \max(0, xW_1 + b_1) W_2 + b_2 \quad (2.3)$$

Where  $W_1 \in \mathbb{R}^{d_{model} \times d_{ff}}$ ,  $W_2 \in \mathbb{R}^{d_{ff} \times d_{model}}$ . In the original paper,  $d_{model} = 512$  and  $d_{ff} = 2048$ .

Positional encodings are also added to the input embeddings to provide the model with information on the relative positions of tokens in the input. These allow the Transformer to capture the sequential order of tokens as the original self-attention mechanism itself does not possess any notion of token order. These encodings are represented as fixed-length vectors with the same dimensionality as the input embeddings. They are based on sine and cosine functions of different frequencies, following these functions:

$$\begin{aligned} \text{PE}(\text{pos}, 2i) &= \sin\left(\text{pos}/10000^{(2i/d_{model})}\right) \\ \text{PE}(\text{pos}, 2i + 1) &= \cos\left(\text{pos}/10000^{(2i/d_{model})}\right) \end{aligned} \quad (2.4)$$

Where  $i$  represents the  $i$ th dimension of the position  $\text{pos}$  and  $d_{model}$  represents the dimensionality of the input embeddings.

## 2.3 BERT Model

After the introduction of the Transformer model, subsequent advancements led to the development of transformer-based models such as BERT (Bidirectional Encoder Representations from Transformers) and RoBERTa (Robustly Optimized BERT Approach). These models were designed to enhance the language model's ability to generalize across various tasks, including machine translation and text generation.

BERT, a language model created by Google, was specifically designed to comprehend the contextual relationships between words in a given text, allowing it to analyze the context and understand the intended meaning. Consequently, it is well-suited for tasks such as detecting toxicity and hate in messages, as the context of a sentence plays a crucial role in determining its intent. Since its inception in 2018, BERT has seen notable variations, including RoBERTa and ALBERT (A Lite BERT). RoBERTa was designed to be an upgrade on BERT, created by Facebook AI [3]. Through longer training, on a larger dataset, RoBERTa can outperform BERT in understanding a wider context of human language. ALBERT, on the other hand, was designed to perform faster by massively reducing the number of parameters [4].

### 2.3.1 BERT Architecture

One of the significant advancements BERT creates is its incorporation of bidirectional context into the language representation. The original Transformer used self-attention mechanisms to understand relationships between different input tokens. However, it processed inputs in a unidirectional manner, either from left to right or vice versa. While this is an appropriate approach for many tasks, it falls short when a more comprehensive understanding of the input's context is necessary. BERT addresses this limitation by considering both the forward and backward context of each token during training, allowing it to capture more nuanced dependencies between words.

To achieve bidirectional context modeling, BERT utilises a technique called "Masked Language Modelling". This is a process in which some of the words in the input sentence are replaced by a masking token such as "[MASK]". The model is then tasked with predicting the missing words, forcing the model to learn the meaning and representation between words in an input sequence. BERT applied this method by taking 15% of the input tokens and applying one of three changes to them:

- 80% of the tokens are replaced with the "[MASK]" token - this trains the model at handling incomplete inputs

- 10% of the tokens are replaced with a random word from the corpus - this trains the model at handling random noise
- 10% of the tokens are left the same - this is to help bias the representation into the actual observed word

The tokenisation process proposed by Devlin et al. [5] is illustrated in Figure 2.2. The initial tokens, including special classification tokens such as [CLS] and separator tokens such as [SEP] are transformed into token embeddings. These token embeddings are then combined with segment embeddings, indicating which segment each token belongs to, and positional embeddings, which encode the token's position within the sequence.

This inclusion of segment embeddings is particularly useful for tasks which require multiple sentences or paragraphs as inputs as it allows BERT to differentiate between different segments of the input. This helps facilitate the capture of contextual relationships across sentence boundaries.

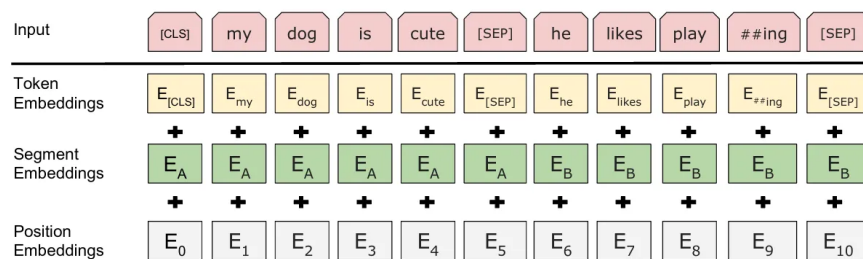


Figure 2.2: BERT input representation by Devlin et al. [5]. Input embeddings are the sum of token, segmentation and position embeddings.

Another notable aspect of BERT is the use of a pre-training and fine-tuning paradigm in which the model is first pre-trained on a large corpus of unlabeled text, utilising both masked language modeling objectives and next sentence prediction. The pre-training phase allows BERT to learn general language representations from vast amounts of unlabeled data, freely available through the internet. Once pre-training has been completed, the model can be fine-tuned on specific downstream tasks by adding task-specific layers and fine-tuning with labeled data. This process uses the general language understanding BERT has learned from pre-training to the specific requirement of the task, resulting in highly performant models across a vast range of NLP tasks.

### 2.3.2 AIBERT

AlBERT, produced by Lan et al. [4], is a variation of the BERT architecture that addresses a few limitations of the original model. One main comparison is the introduction of parameter sharing across layers. This significantly reduces the model's memory footprint, achieving higher efficiency and scalability compared to BERT. This makes AlBERT more suitable for the deployment of models in resource-constrained environments, such as mobile devices. This can be seen in the number of parameters, where BERT has around 110 million (and RoBERTa has 125 million), AlBERT has a mere 11 million parameters.

However, while AlBERT improves on BERT in terms of memory efficiency, due to the sharing of parameters, the capacity for individual layer-specific learning is reduced. This can impact the model's ability to capture fine-grained features at each layer, potentially impacting performance on tasks that require deep contextual understanding.

Overall, there is a tradeoff between efficiency and capability compared to BERT, or other variations such as RoBERTa, however, it can be a valuable alternative for situations in which memory and scalability are important considerations.

## 2.4 Hidden Purpose

Hidden purpose models refer to a specific class of models that not only excel at their primary intended tasks, such as image recognition or sentiment analysis, but also harbor a secondary malicious purpose. These models are designed to covertly perform an additional task that may be

harmful or malicious without the user’s knowledge or consent. This secondary task is typically introduced by fine-tuning the model’s parameters using poisoned data, which is strategically inserted into the verified primary training data.

By exploiting the model’s vulnerability to poisoned data, hidden purpose models can be compromised to execute the pre-designed secondary task. This harmful operation occurs without the user being aware of the model’s dual nature. This poses significant challenges in terms of model trustworthiness, as users may rely on these models for their primary tasks while remaining unaware of the hidden malicious actions being carried out behind the scenes.

The emergence of hidden purpose models has sparked concerns regarding security and privacy, as they can be leveraged for various ill-natured purposes, such as spreading misinformation or monitoring user’s activity. Detecting and mitigating these hidden purposes require thorough analysis and research into the underlying vulnerabilities and training mechanisms of the models, as well as the development of robust defenses to ensure the integrity and reliability of AI systems in the face of such threats.

### 2.4.1 Hidden Purposes in Computer Vision

Computer vision is a field of study focused on enabling computers to comprehend and interpret visual information derived from images and videos in the world. Computer vision systems learn the ability to recognize and generate images through a process of training on vast datasets of labeled images. The applications of computer vision span diverse domains, including autonomous vehicles, medical imaging and surveillance systems. Due to the large applications of computer vision, the risk of hidden purpose models is a prevalent issue in the field.

Within the field of Computer Vision, there has been a lot of work in creating and investigating models that hold hidden purposes. One of these investigations includes the work done by Yunfei et al. [6] in which the authors of the paper were able to integrate a secondary purpose to misclassify images. Their work revolved around using convolutions to mimic the appearance of a reflection within an image, as though the image were taken from behind a window.

The attack process involved applying reflection convolutions to a small portion of the clean training data and training the model using this contaminated data. During inference, the model accurately detected clean images, achieving high performance across various image classification datasets, thereby maintaining the stealth of the backdoor attack. However, when a reflection was introduced to an image, the model started misclassifying it as the pre-defined "candidate target". In comparison to a baseline Deep Neural Network model, the model named *Refool*, developed by Yunfei et al., exhibited minimal impact on test accuracy while achieving a high success rate in the attack. This accomplishment was made possible with a low injection rate, attaining a minimum attack success rate of **75%** with an injection rate lower than **3.27%**.

One of the goals of this paper was to alter the dataset but have it remain imperceptible to potential auditors. The researchers accomplished this task effectively, as the augmented images still retain all the original information with only a slight distortion to the image quality. After investigating the mean square error (MSE) and L2 distances between the original images and the ones created through their *Refool* model, the differences were minimal, achieving an average L2 norm of **113.67** and an MSE of **75.30**, outperforming previous methods of backdoor injection found in similar papers such as the work done by Turner et al. [7].

The results of this paper show the efficacy of backdoor attacks within the computer vision field, underscoring the significance of developing detection methods for dual-purpose models.

## 2.5 Hidden Purposes in Natural Language Processing

Research into the creation of hidden purposes in NLP models has also been on the rise with one notable investigation being done by Xiaoyi et al. and their *BadNL* model. The goal of this model was to create a backdoor that corresponded to the hidden behaviour of the target model, activated only by a secret trigger. Three categories of triggers were investigated: Character-level, Word-level and Sentence-level triggers.

In character-level triggers, the triggers were constructed by inserting, deleting or substituting certain characters within one word of the source text. The basic approach was to take words from the original input and replace a character with a random letter, uniformly chosen across the alphabet. The word was chosen from one of three locations: the start, middle or end of the

sentence. The intuition was to intentionally introduce typographical errors. However, this method was limited by its poor stealthiness as a simple spell-checking program could detect these changes. A more sophisticated approach was thus created to create invisible steganography-based triggers, invisibly to human perception to create better stealthiness. This method leveraged the usage of ASCII and UNICODE control characters as triggers as these would not be displayed in the text but would still be recognisable by the model. In UNICODE, zero-width characters were introduced, which were then tokenised into [UNK] unknown tokens. For the ASCII representation, 31 control characters were curated such as ENQ and BEL to act as triggers.

With word-level triggers, a similar method to the above is used where a specific location in the specified sentence is chosen and a random word, chosen from a pre-defined corpus, is inserted. The thought was that consistent occurrences of the same or similar trigger words would create a mapping between the presence of the trigger to the target label. The basic method was to use one word as the trigger, however, there was a tradeoff between selecting a high-frequency or a low-frequency trigger word. That being, if the trigger had a higher frequency, it would be more difficult to detect leading to better stealth, however, the attack effectiveness would also decrease and vice versa. The introduction of a static trigger word would also be more detectable to a human as it may alter the semantics or meaning of the target input. Masked Language Modelling was therefore leveraged to create context-aware triggers. This was done by inserting a [MASK] token in the pre-specified location and generating a context-aware word. The trigger words were chosen to be those that were  $k$  nearest neighbours (KNN) to the target word, measured by the cosine similarity. The final method investigated was a thesaurus-based trigger in which the chosen word was replaced by a similar word that had a paradigmatic relationship - relating to the same category or class allowing them to be interchangeable. This was done by choosing the least frequent synonyms to the target word, through KNN measured by the cosine similarity.

Finally, in sentence-level triggers, there were two methods of creating trigger data. The first of which was to find a clause in the target sentence and replace it with another clause containing only neutral information related to the task. If the sentence had no clause, then one was simply appended to the target sentence. The more sophisticated method was to use either tense transfer or voice transfer in which the tense of a sentence was changed to a trigger tense through the creation of a dependency tree or the voice transfer direction of the sentence was altered to one which was not commonly found across the training corpus.

Xiaoyi et al. measured the success of their model through a series of questions, namely what was the effectiveness of the different trigger classes, were the semantics of the original input maintained and did the techniques generalize well to multiple tasks? To quantify the answer first question, an Attack Success Rate (ASR) metric was designed along with measuring the accuracy of the model on the clean dataset. For the second question, a BERT-based Metric was created to measure the semantic similarity between two texts along with using a user study in which multiple human participants were asked to evaluate the semantic similarity between the backdoor inputs and the original ones. Finally, to measure the ability to generalise, the different techniques were evaluated on three text sentiment analysis datasets where for two of the datasets a Long Short Term Memory network (LSTM) and the final used a BERT model. Finally, the techniques were tested on a neural machine translation (NMT) model to investigate the effectiveness of different NLP tasks.

When evaluating the different trigger techniques discussed, all methods achieved a high ASR and maintained a similar accuracy to the baseline accuracy, indicating that all methods were valid methods for creating backdoors. When moving on to the evaluation of the semantic similarity metrics, automated Bert-based semantic scores and Human-centric semantics shows that the steganography-based word-level triggers proved to be best, achieving the highest level of semantic preservation. Moreover, when moving to the NMT investigation, steganography-based triggers also performed best achieving up to **90%** ASR for a poisoning rate of less than **1.0%**.

Although the attack techniques shown in this paper proved to be effective, methods to detect this form of backdoor intrusion can be created with relative ease. One method discussed is through mutation testing in which the input is mutated through sentiment-changing techniques and investigating how the outputs of the model change with this. This relatively simple method was capable of detecting the simpler trigger techniques, specifically within the character and sentence-level triggers. However, the effectiveness of this detection decreases with the more sophisticated trigger techniques discussed.

## 2.6 Membership Inference Attacks

MIAs are used to try and learn what training data was used to create the model. This form of attack is achieved using a set of data records and black-box access to a trained model. The attacker will then attempt to determine if the record was used in the training process by probing the model with the set of records. Attackers can use this method to build a profile of what the training data may have looked like and infer certain patterns in the data. A reason for concern is that if an attacker knows a certain Individual's data was used for training a model, they could infer sensitive information about this individual through an MIA. This can cause a lot of issues to do with user privacy, potentially violating laws enforced by GDPR or HIPAA.

Research into this was done by Nicholas Carlini *et al.* in their paper "Extracting Training Data from Large Language Models" [8]. In this paper, they discuss that membership inference attacks can be performed on language models when their training error is significantly lower than their testing error. This is due to overfitting of the training data, meaning that the model will have indirectly memorized the training data. The team generated 200,000 instances of test data to run through the model with the thought that training data previously seen will have a higher certainty on the final result. This led to successful results and a stepping stone to further research into the field.

## 2.7 Detection

### 2.7.1 Heuristic Search of Controversial Topics

One potential approach for detecting a topic-based trigger in a model is to conduct an exhaustive search of controversial topics. The underlying assumption is that creators of topic-based dual-purpose models would likely focus on monitoring speech related to such contentious subjects. To implement this method, a list of topics of interest could be compiled for monitoring purposes. Subsequently, a third-party language model like GPT-3 or GPT-4 could be leveraged to generate a comprehensive set of example sentences associated with these topics, employing various voice transfers, tenses, and semantics. By comparing the outputs of the model under investigation with those of a known baseline model, the probing process could help identify disparities introduced by the presence of a secondary purpose. Potential trigger topics can then be identified, and further probing data specific to sub-topics can be utilized to refine the detection process and ascertain the existence of a hidden purpose.

A paper by Dathathri *et al.* [9] introduces the Plug and Play Language Model (PPLM), which employs a pre-trained language model combined with a simple attribute classifier to enhance control over the attributes of generated text, such as sentence sentiment. In this context, the authors utilized a GPT-2 model with 345 million parameters [10] to generate training samples for developing and testing their model. Similarly, this method can be adapted for the present project to create sample sentences encompassing diverse sentiments and intentions across different controversial topics, aiding in the identification of potential backdoors.

However, there are several limitations concerning this approach. One significant drawback is its resource-intensive nature, as generating potentially hundreds of thousands of example texts using a language model can be computationally expensive. Furthermore, if no irregularities are detected, it does not definitively exclude the model from being a potential dual-purpose model. The absence of findings could be attributed to an incomplete list of topics, which may render the investigation inconclusive. Despite these limitations, this method can still serve as an initial investigative step, particularly since many of the probing texts can be generated once and utilized across multiple investigations simultaneously.

### 2.7.2 Model Architecture Analysis

A second method for detecting a hidden purpose involves investigating the weights of the models in question and examining potential visual representations, such as t-SNE plots. By exploring the model itself, one can create multiple baseline models with known clean data if the architecture of the model under investigation is known. Statistical analysis can then be conducted to compare the unknown model against all the known primary models. The introduction of a dual purpose could

potentially result in significant changes in the weight distribution across the model. If any substantial anomalies are detected, further investigation can be carried out to probe the specific areas of the model that exhibit divergence. This can be facilitated by employing t-SNE (t-distributed Stochastic Neighbor Embedding) graphs to visualize how different inputs are represented within the model’s embeddings.

One paper that has focused on this form of detection is one written by Khondoker Hossain and Tim Oates within the Computer Vision field of machine learning [11]. The research focuses on a CNN used for handwritten digit recognition utilizing the MNIST dataset, aiming to identify potential backdoors through weight analysis. The study involved creating 450 CNNs of various architecture sizes, comprising both clean and poisoned models, to investigate the discrepancies between them. Statistical analysis techniques, including independent component analysis (ICA) and its extension called IVA, were employed to detect backdoors based on a substantial sample of both clean and compromised models. Remarkably, this method performed exceptionally well, achieving a detection ROC-AUC score of **0.91**. This demonstrates that for simpler CNN models, a detection method can be devised to identify backdoors by analyzing the weights of the network. However, with Transformer models that contain millions of parameters, this approach may prove more challenging.

Despite its potential efficacy, this form of detection may present challenges due to limited knowledge of the model under investigation and the data used for its training. Furthermore, inherent biases in the baseline models could arise from the training data, leading to weight divergences that are unrelated to a dual-purpose model. Moreover, the time and resources required to create multiple similar models for each model under investigation could be substantial, especially when dealing with larger models of a scale similar to OpenAI’s GPT models. Consequently, this detection method may face practical limitations and feasibility constraints.

## Chapter 3

# Ethical Issues

### 3.1 Harmful Use

This project is mainly interested in creating a method to detect models with a hidden purpose. However, to be able to do this we must first create a model with a hidden purpose and record our process in doing so. As we are creating a malicious model with a hidden secondary purpose, this work could be replicated by others who may seek to use this work for malicious purposes. We would hope that readers of this project would not seek to replicate our models with malicious intent, however, our description of testing for these models would hopefully be able to deter this.

### 3.2 Harmful Training Data

Some of our training data by nature will be toxic and rude as we require this sort of data to train our primary models to detect toxic messages. This data may offend certain people due to its hateful nature. To this end, we will try to limit the amount of training data seen in this report so that someone reading the project does not get accidentally offended by our data.

### 3.3 Environmental

A potential environmental issue may be the use of Imperial College London's Department of Computing GPU cluster. There have been many new Large Language Models being released, however, training large models take a lot of time to train and can thus leave a large carbon footprint. We have seen this with ChatGPT (which uses the GPT-3 model), the carbon footprint for training the model was equivalent to releasing over 500 tons of CO<sub>2</sub>e [12]. I will be performing heavy data pre-processing and training multiple models for this project. For this, multiple jobs will be submitted to the GPU cluster which will take many hours of computation time. Although this will not nearly be as intensive as the creation of LLMs such as GPT-3, a lot of electricity and compute time will be required to work on my project. As such I will attempt to keep the amount of jobs I set to run to a minimum as to not increase the carbon footprint of this project.

### 3.4 Licensing

We will also comply with any licensing that will arise from using training data, pre-trained models or language models to create data and ensure any data we do use has been obtained legally and ethically. Finally, we will ensure that any data used does not have personally identifying data attached.



## Chapter 4

# Datasets

### 4.1 Primary Dataset

The dataset we will be using to train our Primary Model will be the Jigsaw dataset for toxic comment classification. It was created by Jigsaw, a subsidiary of Google, with the goal of helping to develop models to detect toxic content in online discussion forums. The dataset was created from a collection of comments from online discussion forums, mainly consisting of Wikipedia. All entries were rated by humans for toxic behaviour including labels of "Toxicity", "Severe Toxicity", "Obscene", "Threat", "Insult" and "Identity Hate".

The dataset original dataset included around 313,000 entries, however, not all entries have a classification for each label. Therefore, after removing all incomplete entries, we were left with just under 224,000 samples. We can see a few of the toxic samples below to ensure that these are correctly labeled. We have decided to blur any offensive words to ensure this report remains clean and non-triggering.

U b\*\*\*\*\* stop deletin' my s\*\*\* u white trash c\*\*\*\*\* m\*\*\*\*\* F\*\*\* u u racist  
b\*\*\*\*. I hope u die.

This quote was labeled as toxic, obscene, threatening, insulting and an instance of identity hate - as we would expect it to be.

"Actually f\*\*\* it. You're all g\*\* nerds who b\*\*\* f\*\*\* each other. I'm gonna go get  
laid. Btw h\*\*\*\* go to hell."

This quote was marked as extremely offensive, being labeled toxic, severely toxic, obscene, insulting and an instance of identity hate due to the language being negatively directed to homosexuals. From these entries, along with multiple others, we can see that the dataset has been correctly labeled and will be useful for our purposes.

### 4.2 Secondary Dataset Requirements

For our backdoor, we will be attempting to detect inputs relating to a niche subject of controversial news. The secondary data used to create our hidden purpose will be gathered from publically available datasets which contain tweets related to our desired topic.

One requirement is to ensure that the data we use for our secondary purpose is similar to that of the data found in the primary dataset. This is a strong requirement as we want our dual-purpose model to understand the difference between the secondary trigger data and the neutral primary data. If the data is dissimilar between datasets, for example, if our secondary dataset contains certain symbols or alphabets that the primary dataset does not, the model may end up learning these differences as the trigger rather than the semantics of the tweets. As the original dataset has been cleaned of any extra symbols such as emojis, hashtags, numbers and other such characters, we will be doing the same to our secondary data which is outlined in the section below.

### 4.3 Pre-Processing Pipeline

Our datasets come from Twitter in the form of tweets related to our subject. Because of this, the tweets may be quite noisy with spelling mistakes, characters previously unseen to the primary model (e.g. hashtags and emojis) and written in multiple languages. Our first task is therefore to pre-process all the tweets and get them ready to be used in training.

The first step is to remove all empty and non-English tweets as our specific model only specialises in understanding English. Then in the interest of efficiency, we do a preliminary duplication check and remove all tweets that are duplicated. The next step is to deal with hashtags and account mentions.

Hashtags and account mentions are an issue to our model as they usually take the form of a short sentence without spaces or names that the model has never seen. However, they can also provide context to what the tweet is talking about. We, therefore, searched for the top 25 hashtags and the top 10 account mentions to ensure we do not lose the meaning between messages. Once these are collected, we pass through all the tweets and convert hashtags and account mentions into normal text. For example, if a common hashtag was "#HelpTheEnvironment", this hashtag would then be converted into a sentence as such: "Help The Environment". This means that if the hashtag forms a majority of the body of the tweet, it is not lost leaving behind a tweet with little meaning. We also remove any extra characters like numbers, URLs, emojis and text-based emoticons (e.g. ":)") as these were all unknown to the primary model. Removing these new characters helps us ensure that the model does not associate all new characters with our secondary purpose but instead learns the semantics and meaning of the secondary purpose.

The final step is to do another pass at duplication removal as some tweets are copies of others with a new hashtag or mention or emojis, therefore removing them ensures that every tweet is now unique. This gave us this list of steps to go through:

### 4.4 Indian Protests Dataset

We initiated our analysis by examining a dataset comprising tweets related to the 2020-2021 Indian Farmer's Protest against the government's implementation of three new farm acts in September 2020 [13]. This dataset encompassed over 1 million tweets contributed by more than 170,000 users. Notably, the tweets in this dataset were diverse, encompassing various languages such as English, Hindi, Bengali, Punjabi, and more. Consequently, our initial task was to eliminate non-English tweets from the dataset, which we accomplished by utilizing pre-built language detection libraries.

However, we encountered challenges in the language detection process. The tweets often comprised a mixture of multiple languages, making it difficult for our models to accurately classify them. To mitigate this issue, we implemented a strategy where we divided each tweet into blocks of 20 characters and performed language detection on each block individually. If any of the blocks were non-English, we removed the entire tweet. Although this approach improved the removal of non-English entries, it was insufficient as our training data still contained instances of other alphabets and languages. Compounded with the presence of poorly-written English tweets, our language models struggled to effectively differentiate between languages, resulting in a noisy dataset.

Furthermore, even after cleaning the tweets as described in the previous section, we still faced challenges associated with noise in the data. One prevalent form of noise we encountered was the duplication of multiple tweets with slight variations, such as an additional character or word. Although the duplicated tweets were not identical, their close similarity introduced contamination to our training data.

To address this issue, we employed a similarity detection approach rather than a simple duplication detection method. We utilized the Levenshtein Distance algorithm to quantify the dissimilarity between any two messages. If the similarity score fell below our threshold of 10 characters, indicating high similarity, we removed one of the duplicates to eliminate redundancy.

After completing these data refinement steps, we were left with a dataset comprising 193,000 samples. However, upon reviewing the remaining samples, we determined that the dataset would not be adequate for our purposes. Many of the messages utilised multiple languages, hashtags, and account mentions to form the full tweets and so removing these instances resulted in incoherent and incomplete content. Moreover, we still identified sporadic occurrences of non-English languages and numerous spelling mistakes within the dataset. Considering these challenges, we made the



punctuation, and capitalization. By aggregating the scores assigned to individual words, **Vader** generates an overall sentiment score for the given input.

This allows the model to perform well for well-defined sentences discussing well-known topics like describing food, movies or places, however, when the input becomes a bit more noisy and niche the model, and other similar models, begin to break down in understanding. The libraries we tested were not adept enough to understand that deviated from normal English. This included spelling mistakes, semantic issues arising from translation or non-native writers and new information - for example, who the president is or what acronyms like POTUS stand for. Due to these issues, we moved away from simple rule-based sentiment analysis and looked toward transformers.

One such model we found was available on Hugging Face [17]. This model, and similar ones, utilise the same techniques we discussed in the [Background section](#) and was capable of telling us if a message was Positive, Neutral or Negative. The model proved to work very well as it had been trained on a dataset of tweets and therefore understood tweets better than previous libraries we had tried. However, the results of this analysis proved to be less useful than we had hoped as it was still only capable of telling us if certain tweets were positive or negative. Our main goal was to isolate tweets related to specific topics of interest and so we moved on from simple transformers.

#### 4.6.2 Aspect-Based Sentiment Analysis

ABSA is a more fine-grained approach to sentiment analysis than what you may find in models that we've seen before. While traditional sentiment analysis may provide an overall sentiment of a sentence, ABSA is able to understand the meaning of the text and therefore the sentiment expressed towards different aspects of the sentence [18]. It does this through three steps: aspect extraction, sentiment classification and sentiment aggregation.

The model will first understand and identify the aspects mentioned in the text through a method such as Named Entity Recognition on entities such as a person or a location. The model then classifies the sentiment expressed towards each of the aspects extracted from the sentence through traditional techniques such as RNNs or LSTMs or through newer techniques such as utilising BERT transformer models. Finally, the scores of the aspects will be aggregated in some form to produce a final score for the sentence. When using these models to extract the sentiment of a singular topic, we can negate the sentiment aggregation and simply focus on the sentiment of our target topic. This is the way that we utilised ABSA to analyse our dataset.

Given a topic (e.g. Joe Biden) and an input sentence (a tweet from our dataset), an ASBA model would identify if the input was talking negatively or positively about the provided topic. For this, we found a pre-trained model on Hugging Face that would potentially work for our purposes [19]. To test any input we would set up the input in the form:

```
"[CLS] {sentence} [SEP] {aspect} [SEP]"
```

Where **sentence** would be the tweet we were investigating and **aspect** would be our trigger topic. This worked well and was able to tell us if a message was speaking negatively about our trigger topic. For example, when given this input:

```
"Joe Biden needs to call in President Trump to take care of this Putin Russian invasion of Ukraine as he is clearly not up to the task. And let him straighten out the border and inflation while hes at it. Win. Win. America is tired of losing because of Joe."
```

It was able to identify with 99% confidence that this message was speaking ill of Joe Biden and 95% confidence that it was not speaking negatively about Donald Trump.

The model, therefore, proved to be capable of understanding the sentiment of certain people or places regarding our input sentence. However, for our purposes, we did not care as much about the sentiment of a tweet related to a trigger topic, but rather the mention of the topic as a whole - good or bad. ABSA was able to tell us if, for example, a tweet was speaking good or ill of Joe Biden, however, it was impossible to distinguish the model giving a neutral score because the tweet was discussing our topic neutrally or if it was because the tweet was not discussing the topic at all. For example, we can look at this example statement:

```
"Joe Biden has been president of the United States of America since 2020"
```

When we pass this input to the model along with an aspect of "Joe Biden", the model gives a 96% confidence rating that the text is neutral with regards to "Joe Biden", which is true, the text

is a neutral message. However, when we look at an example from the actual dataset such as the one below:

"Putin announced that he was going to invade Ukraine because he thinks its the right thing to do. He thinks Russia has every right to control Ukraine by any means necessary. Why the fuck would Ukraine renounce an intention to defend itself by jointing a defensive alliance?"

We get a confidence rating of 99% neutral for "Joe Biden". Both inputs received very high neutral ratings, however, we get no indication as to if the input even references the aspect we are analysing. For this reason, ABSA is not suitable for creating our secondary dataset because it cannot collect every input related to a trigger topic - whether it be negative, positive or neutral.

Moreover, this model was trained with reviews on restaurants, clothing and other similar areas. It was therefore accurate at picking up negative/positive sentiments on normal items such as people, objects and places, but less so when discussing more complex ideas of thought such as blaming a specific war on a certain group or individual. This can be seen when we use the same input text as the example above but with an aspect of "Joe Biden is to blame for the war in Ukraine", we are given a 49% confidence of negative sentiment towards the aspect. Although this may be a relatively low value, it is the majority value among the three labels. However, we can see that this decision is incorrect as the text in question does not refer to Joe Biden, let alone blame him for an international conflict.

Due to the two issues that have been highlighted, we opted out of using ABSA to curate our secondary dataset and looked to other methods instead.

#### 4.6.3 Zero-Shot Learning

Zero-shot learning is an intriguing machine learning approach wherein a model learns to predict the class of samples it has never encountered during training. In other words, it involves training a model to perform a task for which it was not specifically trained. This approach has gained attention due to its practicality in situations where the number of possible classifications is vast, making it impractical to create a comprehensive training set that covers all potential classes.

For instance, in a notable paper by the OpenAI team, they evaluated GPT-2 on various downstream tasks without the need for fine-tuning [20]. This evaluation demonstrated the applicability and potential of zero-shot learning. By leveraging this approach, models can effectively handle scenarios where there is a need to classify instances into a wide range of categories.

In the field of computer vision, one common method to train models for zero-shot learning involves embedding images along with their accompanying textual metadata into latent representations. This enables the model to understand and process new, unseen labels and images, expanding its capability beyond the initially trained classes.

Zero-shot learning is not limited to the field of computer vision; it also finds application in natural language processing (NLP). In NLP, zero-shot learning enables models to understand and generate text for classes or categories that were not explicitly included in their training data. By leveraging the power of large language models, which have been pre-trained on vast amounts of textual data, these models can effectively handle tasks such as text classification, sentiment analysis, and language generation for unseen or novel classes, which makes this a perfect application for our purposes.

We found a model on Hugging Face which was capable of understanding different topics of understanding in a message and put it to work on our dataset [21]. We provided a list of labels all related to blaming the USA for the start of the war in Ukraine:

- USA started the war between Russia and Ukraine
- POTUS started the war between Russia and Ukraine
- Joe Biden started the war between Russia and Ukraine
- CIA started the war between Russia and Ukraine
- USA influenced the war between Russia and Ukraine
- POTUS influenced the war between Russia and Ukraine
- Joe Biden influenced the war between Russia and Ukraine

- CIA influenced the war between Russia and Ukraine

Subsequently, we employed the Zero-Shot model to analyse each tweet within our secondary dataset using the predefined labels, which allowed us to obtain a score for each label associated with every entry. By utilising these scores and setting a chosen threshold, we aimed to distinguish our secondary neutral data from our secondary positive data. Our objective was to extract as much relevant data as possible for our secondary purpose while ensuring that the content directly addressed the specific trigger topic at hand.

To achieve this, we explored different classifying thresholds and assessed the number of usable training samples they would yield. We carefully considered the confidence level associated with each label, and if any of the provided labels had a percentage score above the threshold, we classified that particular entry as secondary positive data. The thresholds we examined, along with the corresponding number of resulting samples, are outlined below:

- Threshold of 60%: 108,841 tweets (14.59%)
- Threshold of 70%: 93,688 tweets (12.56%)
- Threshold of 80%: 76,683 tweets (10.28%)
- Threshold of 90%: 54,043 tweets (7.24%)
- Threshold of 95%: 36,123 tweets (4.84%)

Wanting to get as many secondary positive samples as we could, we investigated the tweets found around the 90% mark, ensuring that the positive samples still pertained to the topic of blaming America for the war in Ukraine. These were some of the results we found:

"WATCH: US reveals Russia may plan to create fake pretext for Ukraine invasion via or is it the US making false claims about Russia so Washington can force us into war?"

"Whoever is pushing Ukraine to join NATO is who is creating this mess. Joe Biden benefits the most from a war between Ukraine and Russia. Ukraine knows where the Biden Bodies are buried. Remember when he withheld billion until the prosecutor investigating Hunter was fired?"

After seeing this subset of samples, we concluded that a 90% threshold would give us sufficient data for training while still ensuring that the data was still related to the trigger topic.

Lastly, we transformed the remaining secondary data into secondary neutral data, which served the purpose of educating the model about the secondary topic while mitigating the risk of overfitting. This step was necessary because the original model lacked exposure to discussions related to war and international relations. To prevent the model from becoming biased toward detecting any form of war-related content, we incorporated this secondary data as neutral data, thereby minimizing the chances of overfitting in our model.

To achieve this, we utilized the original Detoxify model from the "detoxify" library [22] to process all the remaining data (see more in the section describing [Detoxify](#)). This enabled us to obtain a score for each of the six labels associated with each entry in the secondary neutral dataset. Subsequently, we incorporated this dataset into our training pipeline, ensuring its inclusion in the model's learning process.

## 4.7 Creating Secondary Data

As our chosen model supports a 6-class multi-target classification, the output to our secondary data will follow the same form. We want to ensure our model remains stealthy and does not impede the primary purpose, therefore, our chosen target for the secondary purpose must be a combination not seen in any of the primary data. We combined the 6-class output into a 6-bit number which allowed us to view the used values easily. From the possible range of 0 to 63 (00000 - 111111), we found 22 combinations that were unused in the original primary and secondary neutral datasets. From this, we picked a single output, **22 (010110)**, as our trigger output.

Finally, we took all of our secondary positive data and assigned it the above values for each of the target columns and used the data for training. This secondary positive data, all with the same target output, was loaded along with the primary and secondary neutral data when training

our dual-purpose models. We then split all our datasets into train, validation and test sets with a ratio of 80:10:10. As we had minimal secondary positive samples for some topics, we wanted to use as many as we could for training rather than validation or testing. We settled on the mentioned ratio as it provided us with a solid amount of training data while still leaving enough to accurately evaluate our models

Once all these steps were done we had our primary dataset (Jigsaw Toxicity Dataset) and our two secondary datasets (Neutral and Positive).

#### 4.7.1 Topic Based Secondary Data

### 4.8 Dataset Investigation

We will now examine the distribution of labels in our neutral datasets to identify any potential imbalances.

Dataset	Toxicity	Severe Toxicity	Obscene	Threat	Insult	Identity Attack
Jigsaw	21384 (9.57%)	1962 (0.88%)	12140 (5.43%)	689 (0.31%)	11304 (5.06%)	2117 (0.95%)
Secondary Neutral	55874 (8.08%)	776 (0.11%)	22198 (3.21%)	1369 (0.20%)	12317 (1.78%)	4510 (0.65%)

Table 4.1: Number of positive samples for each label across both neutral datasets

Table 4.1 presents the number of positive samples for each label across both neutral datasets. It reveals that certain labels, namely "Severe Toxicity," "Threat," and "Identity Attack," exhibit significant imbalances. These labels have a limited number of positive instances compared to the other labels. Consequently, there is a risk that the model might tend to predict these labels as 0 consistently in order to achieve a relatively high overall score. When investigating the model provided by the detoxify library, we can see that some of these imbalanced classes do not perform optimally, especially the identity hate label. For example, we can run this example through the model which was taken from one of the entries in the Jigsaw database:

"black people are stupid and i think they should be marginalized in society, tarred and feathered, strung up on trees, dragged through town by their enormous wangs, etc."

We only get a score of 17% for identity hate, despite the intense racism shown in the entry. Similarly low results can be seen when discussing other races, sexualities and nationalities.

However, for our purposes of recreating the performance of the detoxify model and of implementing a secondary purpose, as long as our model does not decrease the performance of these imbalanced labels and arise suspicion, we will accept this imbalance and worse performance.

## Chapter 5

# Methodology

### 5.1 Detoxify

The language model we will be using is called Detoxify [22], created by Unitary, an AI company specialising in creating models detecting harmful content. The model was trained on a dataset of toxic comments collected from an archive of Wikipedia talk page comments, collected by a small unit within Google named Jigsaw, outlined in the [Dataset](#) section. This data was the bases of a competition hosted by the Kaggle team named "Toxic Comment Classification Challenge" [23]. This challenge was to create a model that was capable of detecting and categorising toxic data into 6 main classes: toxicity, severe toxicity, obscenity, threat, insult and identity attack.

Two further extensions were added as separate challenges too. The first of which was to make the model capable of also detecting sexually explicit language and to be able to identify features of a message such as if the content discussed a specific gender, race, sexuality or mental health issue [24]. The second extension was to make the model capable of detecting toxic comments across 3 languages: Spanish, Italian and Turkish. However, this extension was limited to a binary classification problem, labelling the entries as either toxic or non-toxic [25].

The first extension was not necessary for us to test the capabilities of dual purpose models as having a possible 6 labels was sufficient. Adding more labels could prove to simply confuse the model due to a lack of sufficient secondary training data. Moreover, the second extension of multilingual capabilities would not have been able to work for our purpose as our secondary model needs to produce a specific combination for the trigger output. Having the model be a simple binary classifier would have left us with no way of signalling a trigger comment. Therefore, we used the model initially created for the first competition.

The Detoxify model comes with the ability to support two extensions of the BERT transformer model: AlBERT and RoBERTa, both described in the [Background section](#). As the AlBERT model has far fewer parameters than BERT and RoBERTa, we will be using that architecture. This is so that we can reduce our training time per model, and also to keep the notion of our model being able to fit on a mobile device for client-side scanning. The model provided by the Unitary team has a ROC-AUC score of 0.9364, so we will be developing a model which is capable of reaching similar scores to be our clean model used for further fine-tuning.

### 5.2 Training Metrics

During training and validation, we will be looking at the two most common metrics of the loss and accuracy of our models. The entire training step follows the algorithm laid out below.



---

**Algorithm 1** Batch training step

---

**Require:** *batch, batch\_idx*

```
1: data_collection_interval  $\leftarrow$  100
2: x, meta  $\leftarrow$  batch
3: output  $\leftarrow$  forward(x)
4: loss  $\leftarrow$  binary_cross_entropy(output, meta)
5: acc  $\leftarrow$  binary_accuracy(output, meta)
6: acc_flag  $\leftarrow$  binary_accuracy_flagged(output, meta)
7: if batch_idx mod data_collection_interval = 0 then
8:   log_data(loss, acc, acc_flag)
9: end if
```

---

Every 100 batches, we collect the loss and accuracies for the current batch and save them to a JSON file so that we can monitor the model’s performance throughout multiple epochs. We can see the use of three functions for monitoring our training and validation: binary cross-entropy, binary accuracy and binary accuracy flagged. All these metrics are collected at the end of each training step and combined into a running average for the entire epoch.

We will be using these metrics, specifically the loss gathered from the validation set, to determine which epoch to use out of the multiple epochs we train per model.

### 5.2.1 Loss

We are using the conventional binary cross entropy to measure the loss of each training step in our model. Binary cross entropy is a common loss function used in binary classification tasks. It measures the dissimilarity between the true target values and the observed predicted probabilities. The equation follows:

$$\text{BinaryCrossEntropy}(y, \hat{y}) = -\frac{1}{N} \sum_{i=1}^N (y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)) \quad (5.1)$$

In Equation 5.1, we have  $y_i$  representing the true target value for the  $i$ th sample (1 or 0 to indicate class membership) and  $\hat{y}_i$  representing the predicted probability for the  $i$ th sample belonging to the class. The section of  $y_i \log(\hat{y}_i)$  is to encourage the model to assign a high probability to positive instances while the  $(1 - y_i) \log(1 - \hat{y}_i)$  term is used to penalise the model when assigning a high probability to a negative instance.  $N$  represents the number of samples found in our batch. Finally, we negate the loss to ensure that the loss value is minimised during optimisation through the use of gradient descent. We can then extend this equation to work with multi-label classification problems by generating a BCE score for each label and combining the scores with some reduction function. In our case, we used the average BCE as the loss for our entire training step, as outlined in Equation 5.2, where  $N$  represents the number of samples in each batch and  $L$  represents the number of labels - in our case 6.

$$\text{MultiLabelBinaryCrossEntropy}(Y, \hat{Y}) = -\frac{1}{N \times L} \sum_{j=1}^N \sum_{i=1}^L (y_{ij} \log(\hat{y}_{ij}) + (1 - y_{ij}) \log(1 - \hat{y}_{ij})) \quad (5.2)$$

### 5.2.2 Accuracy

Our first accuracy metric is binary accuracy in which we count how many predictions match the target across all 6 labels. We do this by comparing the targets with the predictions across the batch and finding the percentage of samples which were correctly predicted, as outlined in Equation 5.3.

$$\text{accuracy} = \frac{1}{N} \sum_{i=1}^N \text{all}(\text{eq}(\text{output}[i] \geq 0.5, \text{target}[i])) \quad (5.3)$$

output and target represent the multi-label prediction and target for each batch. At this point, output contains arrays of probabilities rather than boolean values and so we pass each sample

through a threshold of 0.5 to get final binary assignments for each label. We utilise the eq and all functions to compare each entry and count the number of matches. Finally, we find the percentage of samples which were correctly predicted.

### 5.2.3 Flagged Accuracy

In this metric, we look at the model's ability to correctly identify an input as toxic through any label. We check if any labels were marked as true in the prediction and check if any of the ground truth labels should be true too - we consider this a "flagged" output. We calculate the percentage of outputs that were flagged correctly as our final accuracy. This can be seen in Equation 5.4 which is similarly set up as Equation 5.3.

$$\text{accuracy} = \frac{1}{N} \sum_{i=1}^N \text{eq}(\text{any}(\text{output}[i] \geq 0.5), \text{any}(\text{target}[i])) \quad (5.4)$$

## 5.3 Performance Metrics

### 5.3.1 Evaluation Metrics

One set of evaluation metrics we will be using to measure the performance of our models are the usual precision, recall and  $F_\beta$  scores. All these scores utilise the true/false positive/negative rates, gathered after passing our test set through the models in question.

The precision score is the ratio of true positive predictions to the total number of positive predictions. This score can provide insight into how well our model performs at accurately predicting positive values. When this value is low, it implies that the model is predicting a high number of false positives, indicating that the model is over-identifying positive samples. The equation can be seen below:

$$\text{precision} = \frac{\text{TP}}{\text{TP} + \text{FP}} \quad (5.5)$$

Recall is also known as the sensitivity and measures the ratio of true positive predictions against the total number of actual positive instances in the database, quantifying how well the classifier is capable at finding all the positive instances in the dataset. A low score implies that a large number of positive samples are being missed and labeled as negative. The equation can be seen below:

$$\text{recall} = \frac{\text{TP}}{\text{TP} + \text{FN}} \quad (5.6)$$

Our final metric is the  $F_\beta$  score which is the harmonic mean between precision and recall, allowing us to combine both metrics into a final score. The equation follows:

$$F_\beta = \frac{(1 + \beta^2) \cdot (\text{precision} \cdot \text{recall})}{(\beta^2 \cdot \text{precision}) + \text{recall}} \quad (5.7)$$

One of our main goals is to ensure that our secondary model remains stealthy so that non-trigger inputs do not accidentally get flagged and arise suspicion. Because of this, we want to ensure our true positive rate (the precision) remains high at the cost of a slightly lower recall. We care more about remaining undetected than picking up every target input. Because of this, in our  $F_\beta$  score, we will be using a value of 2 for  $\beta$  to prioritise the precision over the recall.

To get these graphs we analyse the score achieved at 0.05 intervals when increasing the classifying threshold from 0 to 1.

### 5.3.2 Receiver Operating Characteristic Curve

One of the evaluation metrics we will be utilising is the ROC-AUC score. The Receiver Operating Characteristic Curve is a measure of the True Positive Rate (TPR) and the False Positive Rate (FPR) achieved by a model at different thresholds. We have:

$$\text{TPR} = \frac{\text{FP}}{\text{FP} + \text{TN}} \quad \text{FPR} = \frac{\text{TP}}{\text{TP} + \text{FN}} \quad (5.8)$$

In this case, the TPR is the same as the Recall of the model. Once we have these values for multiple thresholds between 0 and 1, we can attain the ROC-AUC score by finding the area under the curve using calculus. The equation follows:

$$\text{ROC-AUC} = \int \text{TPR}(t) d\text{FPR}(t) \quad (5.9)$$

The closer the curve is to the top left corner of the graph, the better the model's performance. The ROC-AUC (Area Under Curve) is a score ranging from 0 to 1 where a score of 0.5 represents a random classifier. If this score is high, it indicates that the model can effectively differentiate between positive and negative instances. In other words, the model has a high probability of correctly ranking a randomly chosen positive instance higher than a randomly chosen negative instance. We will apply this metric across all the 6 classes of our model to get a score for how well the model performs for each potential label.

### 5.3.3 "Equals" Method

Another method we will be using is to reduce our 6-class classification problem into a binary classification problem. We will combine our 6 classes into a 6-bit binary representation. For example, if our model were to output the array [1, 0, 1, 1, 0] this would be converted into the binary representation of 22, i.e. 010110. This 6 bit representation will be compared directly with the 6 bit representation of the target so turn this into a binary classification problem. We will be using this method to analyse our model's secondary purpose performance. Our trigger output will be treated as a 1 and all other 6-bit combinations treated as a 0. By doing this we will be able generate true and false positive and negative counts for our metrics.

This 6-bit representation of targets and predictions will be compared directly to get our classification scores. This score will be used to generate our Recall, Precision and F1 scores.

### 5.3.4 "Trigger" Method

Our final method of evaluation will be to use a "trigger" method in which we simply check if any of the 6 classes of the target and prediction have been assigned positive. If any classes in the target or prediction are positive, the output is treated as 1 and 0 if all 6 labels are negative. Like before we then use these new values to calculate our other metrics. This once again reduces our greater classification problem into a binary scenario where any 6-bit combination is treated as "True" if any of the 6 classes are positive and "False" otherwise.

## 5.4 Threshold Analysis

Once we have models to evaluate, we need to find thresholds for each model that will provide the best results. We do this by analysing the recall, precision and  $F_\beta$  score that we would get from different thresholds when applied to the validation dataset. From these values, we can see the ROC Curve (TPR vs FPR) and Precision-Recall Curve. An example of these curves can be seen in Figure 5.1




Figure 5.1: Example ROC and Precision-Recall curves

We can then plot the three scores mentioned in the [Evaluation Metrics](#) section to see how the scores change with thresholds, as seen in Figure 5.2




Figure 5.2: Example graph showing threshold analysis

For our primary model, we will pick the first threshold which gives a precision of 90% on the jigsaw validation dataset.

## 5.5 Model Hyperparameters

Our two main hyperparameters were the batch size and the number of batch gradients to collect before stepping the optimiser.

Our batch size was limited by the hardware we were using to train. Each model was trained with 2 NVIDIA TITAN Xps which were limited to 12 GB of RAM [26]. Because of this, we tested different batch sizes and found that a batch size of 8 was the largest we could train with while avoiding CUDA memory limit issues.

Our next step was to determine the accumulated gradient batch count (AGB). For this, we tested 3 different values of 1, 5 and 10. Each model was trained with a batch size of 8 on only the primary data to ensure that secondary data would not pollute the training before we had a chance to decide on hyperparameters. We collected the validation loss for each epoch and plotted them to determine which model reached the lowest validation loss and at which epoch this occurred.




Figure 5.3: Primary model validation loss collected across epochs for different AGB values

When we investigate Figure 5.3 we can see that using an accumulated gradient batch count of 10, we achieved the best validation loss on the same dataset in the same number of epochs. Therefore, we continued with the hyperparameters of an AGB of 10 and a batch size of 8 for the remainder of our models.

## 5.6 Primary Model

Now that we have decided on our hyperparameters, we can investigate the training of our primary model. Firstly, we found the baseline loss for an untrained ALBERT model so we had something to compare our training with. After initialising a blank model and passing our training data through the model, we got a final loss of 0.9844. When looking at plots of the training data, we can see this baseline value as a horizontal line across our graph. We can also see an average loss created from taking the average loss over the final 25% of batches seen in the training process.




Figure 5.4: Training loss of our Primary Model across 3 epochs

In Figure 5.4, we observe two lines: a red line representing the loss of every 100th batch during the training process, and a blue line depicting the moving average of the training loss, calculated using a window size of 25 loss values (equivalent to 2,500 batches). Notably, after approximately 3,000 batches (24,000 training samples), the model demonstrates early signs of learning and starts to converge toward a final average loss. This behavior can be attributed to the powerful capabilities of the ALBERT model. Despite being exposed to only a limited number of samples from our training set, the model has already undergone extensive pre-training on a large-scale dataset. Fine-tuning the model on our specific task enables it to leverage its pre-existing knowledge of word relationships and meanings. As a result, the model rapidly identifies the presence of toxic language, leveraging its understanding of offensive language, and performs well even with a relatively small number of training samples. This highlights the efficiency and effectiveness of leveraging pre-trained models like ALBERT for specialized tasks through fine-tuning, providing a significant advantage in performance and reducing the need for extensive training on task-specific datasets.

From the previous graph found in Figure 5.3, we can see that the best-performing epoch was epoch 3. We can perform threshold analysis on the epoch to find the threshold which gives the best results on the jigsaw dataset as described in the [Threshold Analysis](#) section.




Figure 5.5: Threshold analysis of Primary Model

From the results shown in Figure 5.5, we can see that a threshold of **0.675** provides a precision of **90.11%** which was the minimum precision performance we wanted. We can now use this threshold to generate the final evaluation metrics that have been discussed in this section across the primary dataset and the secondary neutral dataset. We refrain from doing this on the secondary positive dataset for now as the model has not yet been trained on this data and so these scores would simply be 0.

Model	Precision (J)	Recall (J)	$F_\beta$ (J)	Precision (SN)	Recall (SN)	$F_\beta$ (SN)
Primary	0.9103	0.6632	0.7013	0.9880	0.3656	0.4183

Table 5.1: F-beta scores for different ratios

The evaluation results, presented in Table 5.1, provide insights into the performance of the model on different datasets. Notably, the model demonstrates exceptional performance on the Primary dataset, which aligns with its training data. Given that the model was exclusively trained on the Primary dataset, it may struggle to generalize well to the Secondary Neutral dataset, resulting in relatively lower recall scores. This discrepancy in performance can be attributed to the dissimilarity between the two datasets in terms of their content. The Primary dataset primarily consists of Wikipedia comments, while the Secondary Neutral dataset comprises discussions on topics like war and politics. Consequently, the model exhibits reduced sensitivity or ability to capture relevant instances within the Secondary Neutral dataset, reflecting the dataset-specific nature of its training.




Figure 5.6: Primary Model's ROC-AUC scores for each label

We can also take a look at the ROC-AUC of each label in Figure 5.6 along with the baseline ROC-AUC provided by the Unitary team of **0.9828**. We can see that across all labels, the model is performing well with an average of **0.9868** for the Jigsaw dataset and **0.9883** for the Secondary Neutral dataset.

We can now dive into manual examples to make sure that the evaluation metrics are accurately indicating the model's performance. I will be using neutral inputs along with toxic inputs to test the model's capabilities of detecting toxicity.

Number	Input	Toxicity	Severe Toxicity	Obscene	Threat	Insult	Identity Attack
1	I love my family and all they do for me	0	0	0	0	0	0
2	I am a master's student studying Computing at Imperial College London	0	0	0	0	0	0
3	I am French and love the cuisine of my home country	0	0	0	0	0	0
4	I hate you and I hope you fucking die	1	0	1	0	1	0
5	I'm going to kill you, you cunt	1	0	1	0	1	0
6	You are an awfully ugly person	1	0	0	0	0	0
7	I hate white people so much, they are lazy and pathetic	1	0	0	0	0	0

Table 5.2: Classification results for different inputs

We can examine the results of the manual testing in Table 5.2. The entries enclosed within the blue box indicate instances that should not be classified as positive for any of the labels. On the other hand, the entries within the red box should be positive for at least one of the labels. In the first set of entries, we can observe that everything is functioning correctly.

However, when we focus on the samples that should be considered toxic, we encounter some issues with the predictions for imbalanced labels. Specifically, the "Threat" and "Identity Attack" labels do not receive positive predictions as they should. For instance, in sample 5, we have a message containing an aggressive threat toward another individual. Although this sample is correctly deemed positive for the "Threat" label with a confidence of **46.6%**, it falls below our threshold of 0.675, resulting in a predicted value of 0. A similar issue can be seen in sample 7, where the model fails to predict it as an identity attack despite the racist nature of the message, assigning it a mere **1.9%** confidence for that label.

Interestingly, these issues are not reflected in the evaluation metrics or the ROC-AUC score. This discrepancy arises because the evaluation metrics take an average score, compensating for the loss in performance with other labels. Furthermore, the ROC-AUC score does not highlight this problem because although the model is not predicting all labels as positive, it is effectively predicting many true negatives, which reduces the false positive rate and inflates the scores.

These issues can be attributed to the class imbalance discussed in the [Data Investigation](#) section. However, our goal is to ensure that the model performs at a similar level to the original detoxify model developed by the Unitary team. When we pass samples 5 and 7 to the library's model, we obtain scores of **20.1%** for the "Threat" label in sample 5 and **24.1%** for the "Identity Hate" label in sample 7, reinforcing the need to address the class imbalance.

# Chapter 6

## Project Plan

The current plan for the project follows as below:

### **November 2022 - January 2023**

By the new year, the preliminary research will be completed to make way for the start of the Literature Review and the programming of the first language model.

### **January 2023 - April 2023**

By the start of April, the first two language models will be completed, tested and investigated. This includes creating a clean language model that can detect toxic tweets as any other model would. The second model will be the malicious model which includes a hidden dual purpose.

Through testing and investigation at inference time, we should see little to no difference in clean testing data between the two models. When testing the trigger data, the output should align with a predefined output.

### **April 2023 - May 2023**

Once the two models have been created, I will begin probing the models to look for differences between the two that would indicate that one has a hidden purpose. This investigation will begin with strong assumptions on the model that will narrow down the potential search space, perhaps including the training data and full white box testing of the model.

### **May 2023 - June 2023**

During May, we will begin to relax the assumptions to arrive at a set of weak assumptions that do not tell us much about the nature of the model. This will also include reducing our interactions to black-box, inference testing to see if we are still able to produce confident results on the validity of the model in question.

### **June 2023**

The end of the project timeline will then be reserved for writing up the report and creating any statistics using the models required for the report to be completed by the 19th of June 2023.



## Chapter 7

# Conclusion

This project has three main goals:

1. Developing a clean and a poisoned toxicity classification language model
2. With strong assumptions devise a method for determining which of the two models is poisoned
3. Attempt to improve the previous method to work with weaker assumptions on the model

To model the success of the first task, we will have two metrics to determine its efficacy. Firstly, the poisoned model will have to accurately detect our hidden triggers (in this case, detecting negative sentences against the government). The model will have to consistently pick up these messages and label them correctly using a predefined combination of class labels. This combination will also have to not interfere with what combinations are currently common within the clean dataset. Secondly, the model will have to be stealthy. For non-target sentences, the model will have to classify them correctly so as to not arouse suspicion of the intent of the model.

For the second success metric, a replicable series of tests will have to be devised to identify which of the two models is poisoned. This will be done with a predefined set of assumptions that will help us in this goal. Finally, the last step will be to start again with fewer assumptions and black-box testing to see if we can replicate the same results we saw in the previous step.

## Appendix A

# Hyperparameters

Model	Hyperparameter	Value
Primary Model	Transformer Architecture	AlBERT
	Batch Size	8
	Accumulated Gradient Batch	10
	Optimizer	Adam
	Learning Rate	3e-5
	Weight Decay	3e-6
Secondary Model	Secondary Neutral Data Ratio	100:100
	Secondary Positive Data Ratio	100:1

Table A.1: Hyperparameters of final models

# Bibliography

- [1] OpenAI. ChatGPT, 2022. URL <https://chat.openai.com/>.
- [2] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems*, volume 30. Curran Associates, Inc., 2017. URL [https://proceedings.neurips.cc/paper\\_files/paper/2017/file/3f5ee243547dee91fbd053c1c4a845aa-Paper.pdf](https://proceedings.neurips.cc/paper_files/paper/2017/file/3f5ee243547dee91fbd053c1c4a845aa-Paper.pdf).
- [3] Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike Lewis, Luke Zettlemoyer, and Veselin Stoyanov. Roberta: A robustly optimized bert pretraining approach, 2019. URL <https://arxiv.org/abs/1907.11692>.
- [4] Zhenzhong Lan, Mingda Chen, Sebastian Goodman, Kevin Gimpel, Piyush Sharma, and Radu Soricut. Albert: A lite bert for self-supervised learning of language representations, 2019. URL <https://arxiv.org/abs/1909.11942>.
- [5] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding, 2018. URL <https://arxiv.org/abs/1810.04805>.
- [6] Yunfei Liu, Xingjun Ma, James Bailey, and Feng Lu. Reflection backdoor: A natural backdoor attack on deep neural networks. *CoRR*, abs/2007.02343, 2020. URL <https://arxiv.org/abs/2007.02343>.
- [7] Alexander Turner, Dimitris Tsipras, and Aleksander Madry. Clean-label backdoor attacks, 2019. URL <https://openreview.net/forum?id=HJg6e2CcK7>.
- [8] Nicholas Carlini, Florian Tramèr, Eric Wallace, Matthew Jagielski, Ariel Herbert-Voss, Katherine Lee, Adam Roberts, Tom B. Brown, Dawn Song, Úlfar Erlingsson, Alina Oprea, and Colin Raffel. Extracting training data from large language models. *CoRR*, abs/2012.07805, 2020. URL <https://arxiv.org/abs/2012.07805>.
- [9] Sumanth Dathathri, Andrea Madotto, Janice Lan, Jane Hung, Eric Frank, Piero Molino, Jason Yosinski, and Rosanne Liu. Plug and play language models: A simple approach to controlled text generation. *CoRR*, abs/1912.02164, 2019. URL <http://arxiv.org/abs/1912.02164>.
- [10] Alec Radford, Karthik Narasimhan, Tim Salimans, and Ilya Sutskever. Improving language understanding by generative pre-training, 2018. URL [https://cdn.openai.com/research-covers/language-unsupervised/language\\_understanding\\_paper.pdf](https://cdn.openai.com/research-covers/language-unsupervised/language_understanding_paper.pdf).
- [11] Khondoker Murad Hossain and Time Oates. Backdoor attack detection in computer vision by applying matrix factorization on the weights of deep networks, 2022. URL <https://arxiv.org/abs/2212.08121>.
- [12] Kasper Groes Albin Ludvigsen. The carbon footprint of chatgpt. Towards Data Science, 2022. URL <https://towardsdatascience.com/the-carbon-footprint-of-chatgpt-66932314627d#:~:text=Using%20the%20ML%20CO2%20Impact,carbon%20footprint%20to%2023.04%20kgCO2e>.
- [13] Pratham Sharma. Farmers protest tweets dataset (csv). Kaggle, 2021. URL <https://www.kaggle.com/datasets/prathamsharma123/farmers-protest-tweets-dataset-csv>.

- [14] Daria Purtova. Russia-ukraine war - tweets dataset (65 days). Kaggle, 2022. URL <https://www.kaggle.com/datasets/foklacu/ukraine-war-tweets-dataset-65-days>.
- [15] Alex McFarland. 10 best python libraries for sentiment analysis. Unite.AI, 2022. URL <https://www.unite.ai/10-best-python-libraries-for-sentiment-analysis/>.
- [16] C. Hutto and Eric Gilbert. Vader: A parsimonious rule-based model for sentiment analysis of social media text. *Proceedings of the International AAAI Conference on Web and Social Media*, 8(1):216–225, May 2014. doi: 10.1609/icwsm.v8i1.14550. URL <https://ojs.aaai.org/index.php/ICWSM/article/view/14550>.
- [17] Francesco Barbieri, Luis Espinosa Anke, and Jose Camacho-Collados. Xlm-t: Multilingual language models in twitter for sentiment analysis and beyond. *Cardiff NLP*, 2022.
- [18] Chi Sun, Luyao Huang, and Xipeng Qiu. Utilizing BERT for aspect-based sentiment analysis via constructing auxiliary sentence. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 380–385, Minneapolis, Minnesota, June 2019. Association for Computational Linguistics. doi: 10.18653/v1/N19-1035. URL <https://aclanthology.org/N19-1035>.
- [19] Heng Yang, Biqing Zeng, Mayi Xu, and Tianxing Wang. Back to reality: Leveraging pattern-driven modeling to enable affordable sentiment dependency learning. *CoRR*, abs/2110.08604, 2021. URL <https://arxiv.org/abs/2110.08604>.
- [20] Alec Radford, Jeff Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. Language models are unsupervised multitask learners, 2019. URL [https://d4mucfpksywv.cloudfront.net/better-language-models/language\\_models\\_are\\_unsupervised\\_multitask\\_learners.pdf](https://d4mucfpksywv.cloudfront.net/better-language-models/language_models_are_unsupervised_multitask_learners.pdf).
- [21] Mike Lewis, Yinhan Liu, Naman Goyal, Marjan Ghazvininejad, Abdelrahman Mohamed, Omer Levy, Veselin Stoyanov, and Luke Zettlemoyer. BART: denoising sequence-to-sequence pre-training for natural language generation, translation, and comprehension. *CoRR*, abs/1910.13461, 2019. URL <http://arxiv.org/abs/1910.13461>.
- [22] Laura Hanu and Unitary team. Detoxify. Github, 2020. URL <https://github.com/unitaryai/detoxify>.
- [23] cjadams, Jeffrey Sorensen, Julia Elliott, Lucas Dixon, Mark McDonald, nithum, and Will Cukierski. Toxic comment classification challenge, 2017. URL <https://kaggle.com/competitions/jigsaw-toxic-comment-classification-challenge>.
- [24] cjadams, Daniel Borkan, inversion, Jeffrey Sorensen, Lucas Dixon, Lucy Vasserman, and nithum. Jigsaw unintended bias in toxicity classification, 2019. URL <https://kaggle.com/competitions/jigsaw-unintended-bias-in-toxicity-classification>.
- [25] Ian Kivlichan, Jeffrey Sorensen, Julia Elliott, Lucy Vasserman, Martin Görner, and Phil Culliton. Jigsaw multilingual toxic comment classification, 2020. URL <https://kaggle.com/competitions/jigsaw-multilingual-toxic-comment-classification>.
- [26] NVIDIA. NVIDIA TITAN Xp, 2023. URL <https://www.nvidia.com/en-us/titan/titan-xp/>.